



# MiningMath Associates


## Dense Cholesky CUDA implementation




**Fabrício Ceolin e Alex Teixeira**




# Block Model




# The current state of the mine planning technology




**Pre-assigned destinations  
Single block values**



**Lerchs-Grossman  
Max. Undisc. Cash Flow**



**Period-by-period scheduling  
with blending**




Taken from Micromine




**Stockpiling and Cutoff policy**

# Mine planning using SimSched Direct Block Scheduler




Leite and Dimitrakopoulos (2009)

**Block values for each model  
and possible destination**



**Indicators with risk profiles**



**Multi-period scheduling  
under uncertainty with  
blending and stockpiling**

8163.66

8165.18

8166.71

8168.23

8169.75

(927,381,86) = 30.374.082 milhões de blocos

848.61

847.90

847.18

846.47

845.75

1355.12

1354.4

1353.69

1352.96

1352.26

Results

Up

North

East

8313.48

8315.00

8316.53

8318.05

8319.57

Scale 1 : 188.022

0

1/4

1/2

3/4

1



-717.56  
-722.87  
-728.18  
-733.49  
-738.80

1406.17 1417.67 1429.18 1440.68 1452.19


2510.74 2522.24 2533.75 2545.25 2556.76

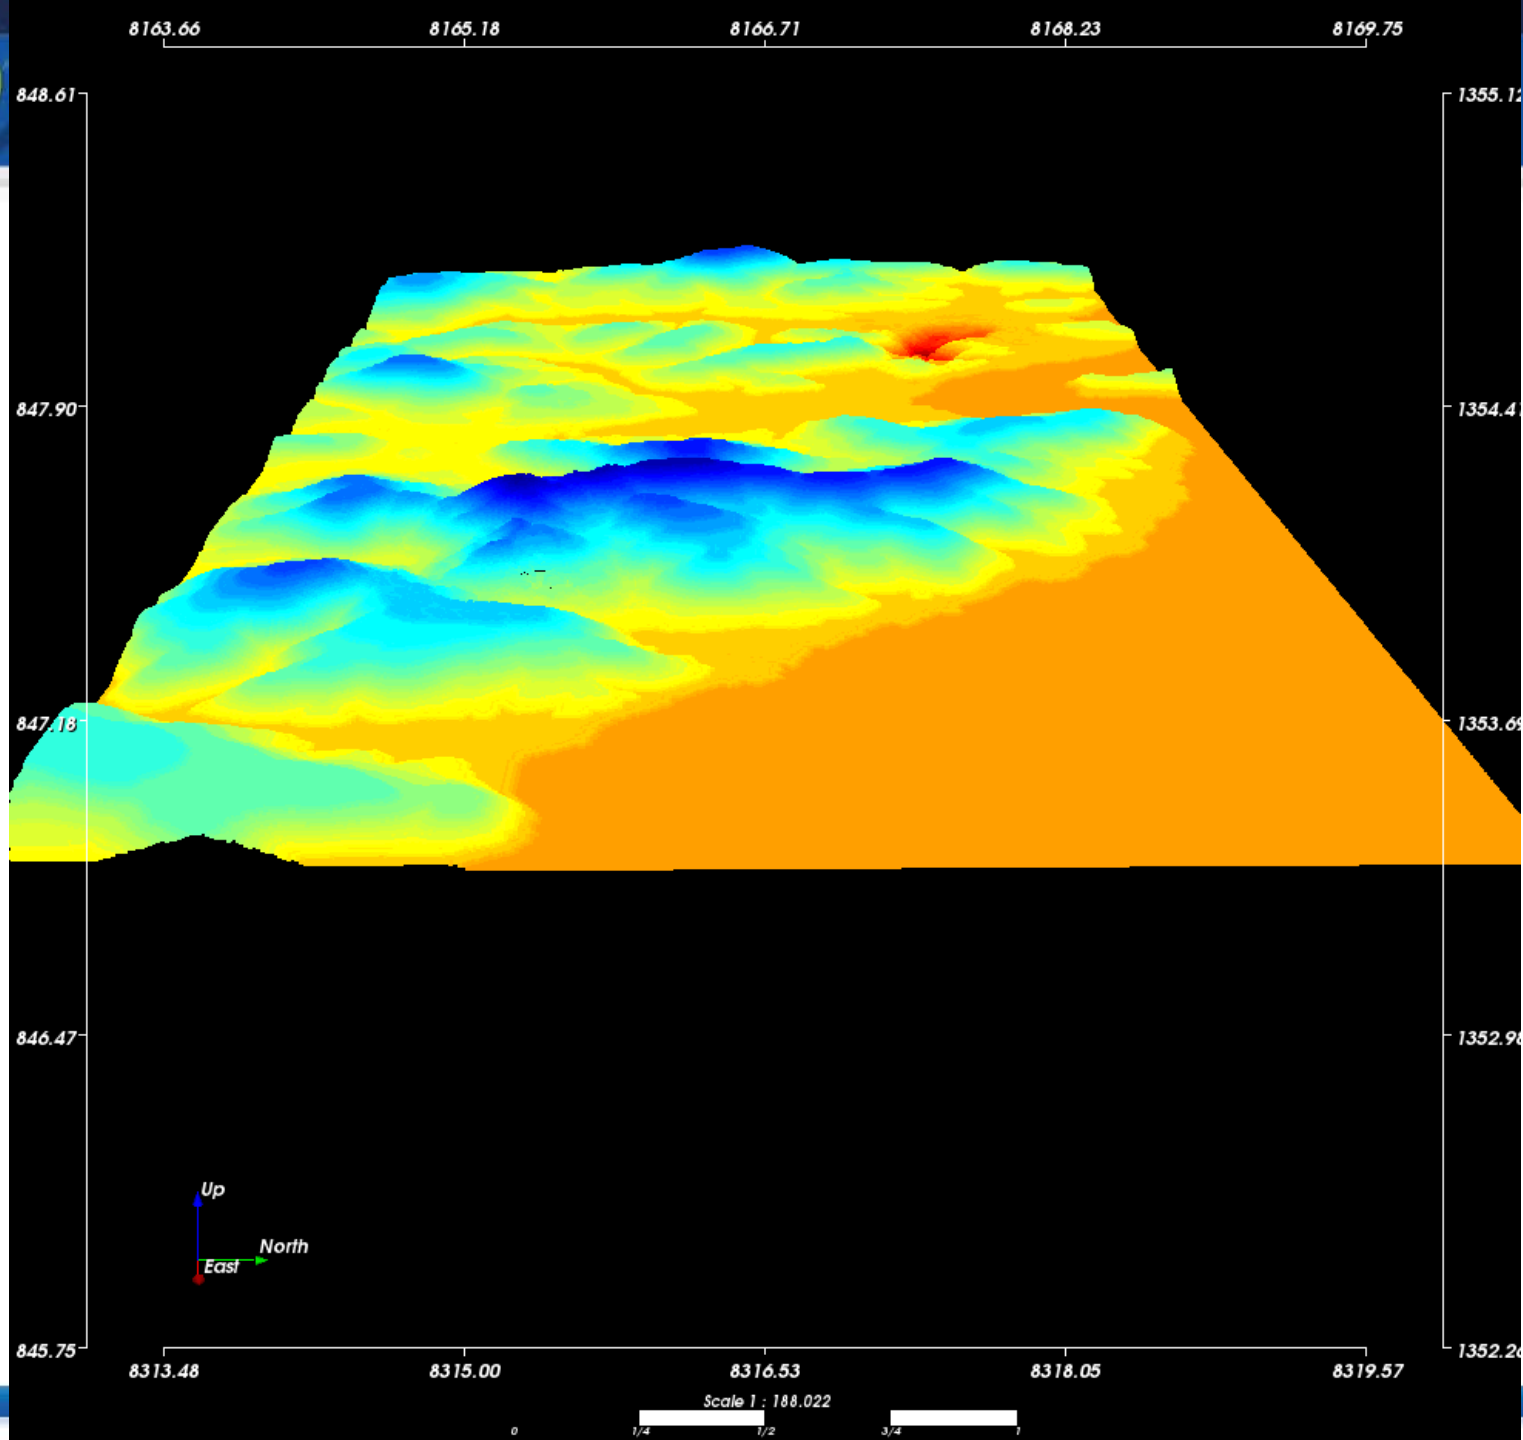
3104.9  
3099.6  
3094.2  
3088.9  
3083.6

Up  
East North


Scale 1 : 1418.94

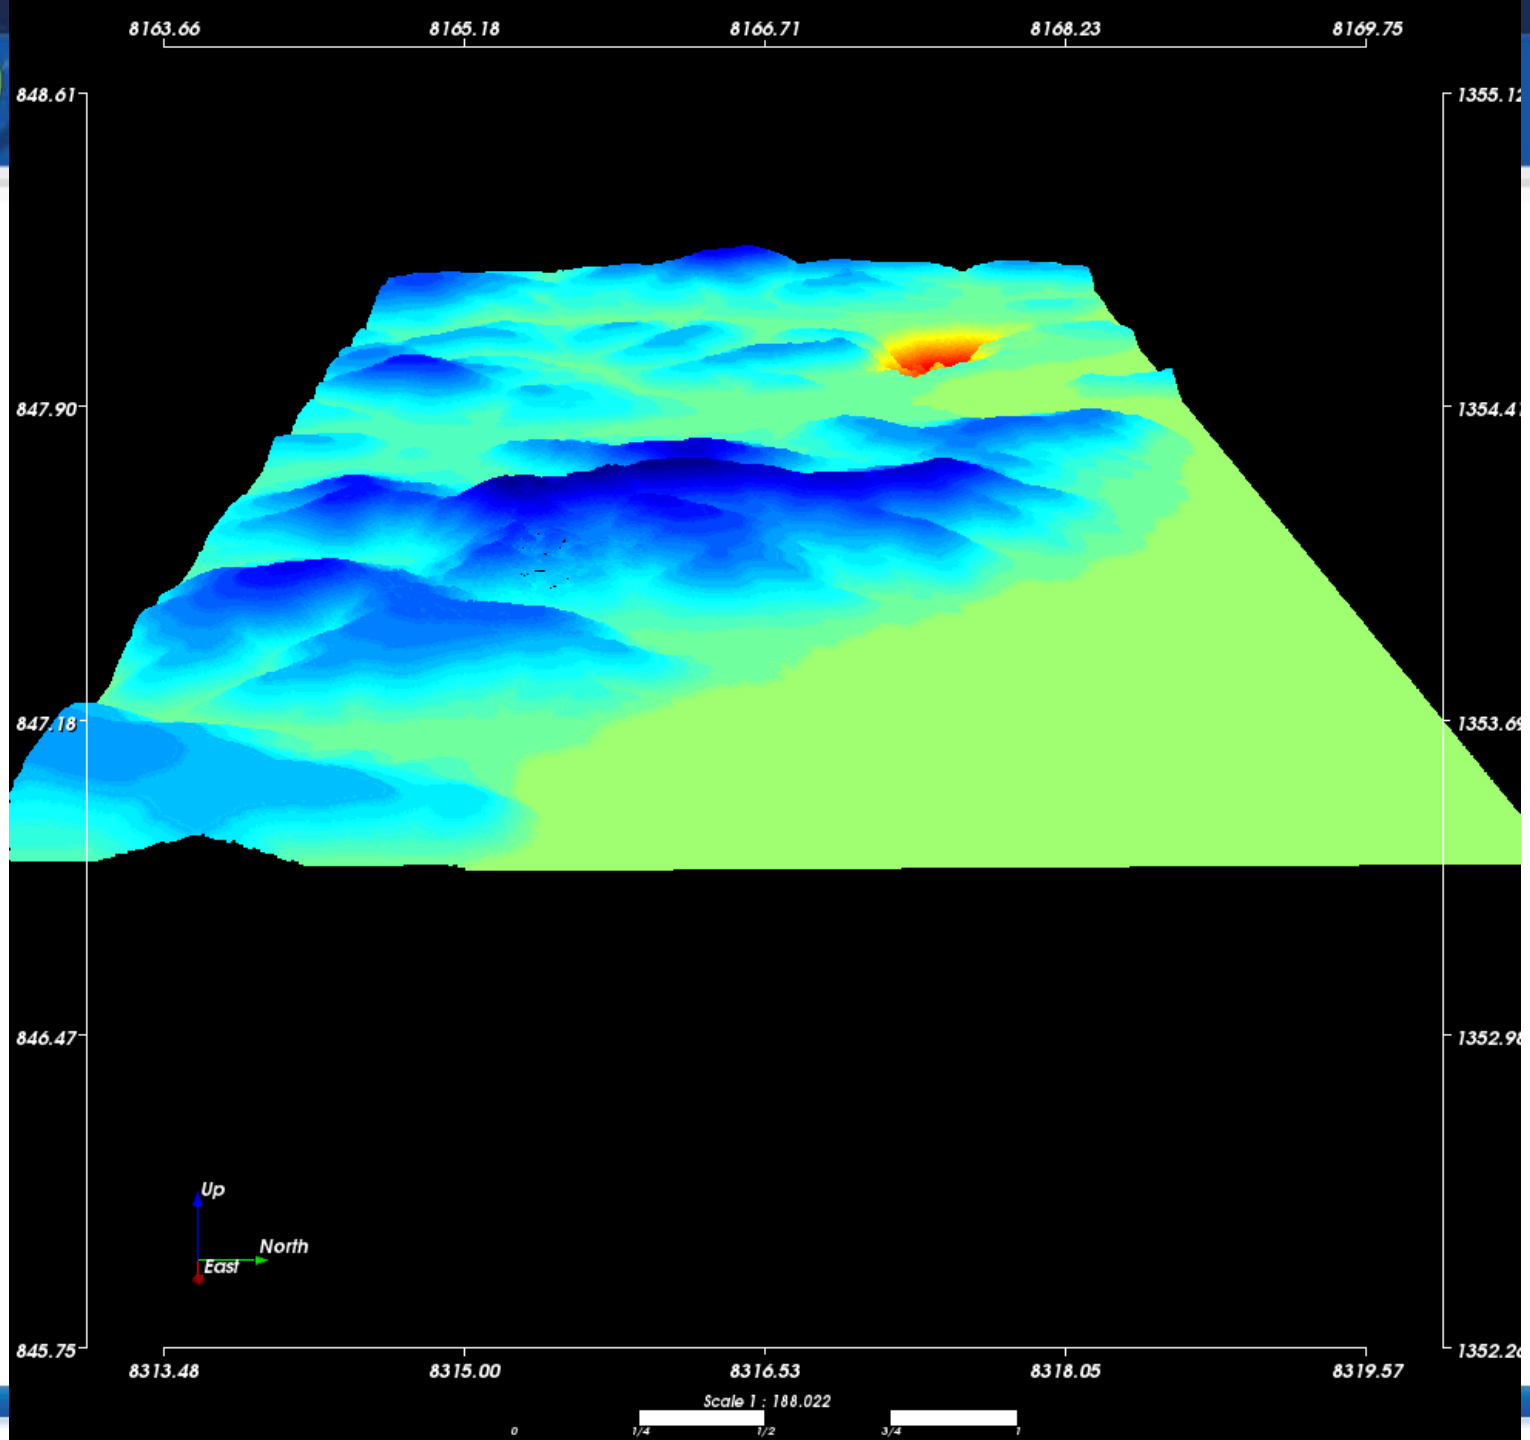
0 1/4 1/2 3/4 1






























































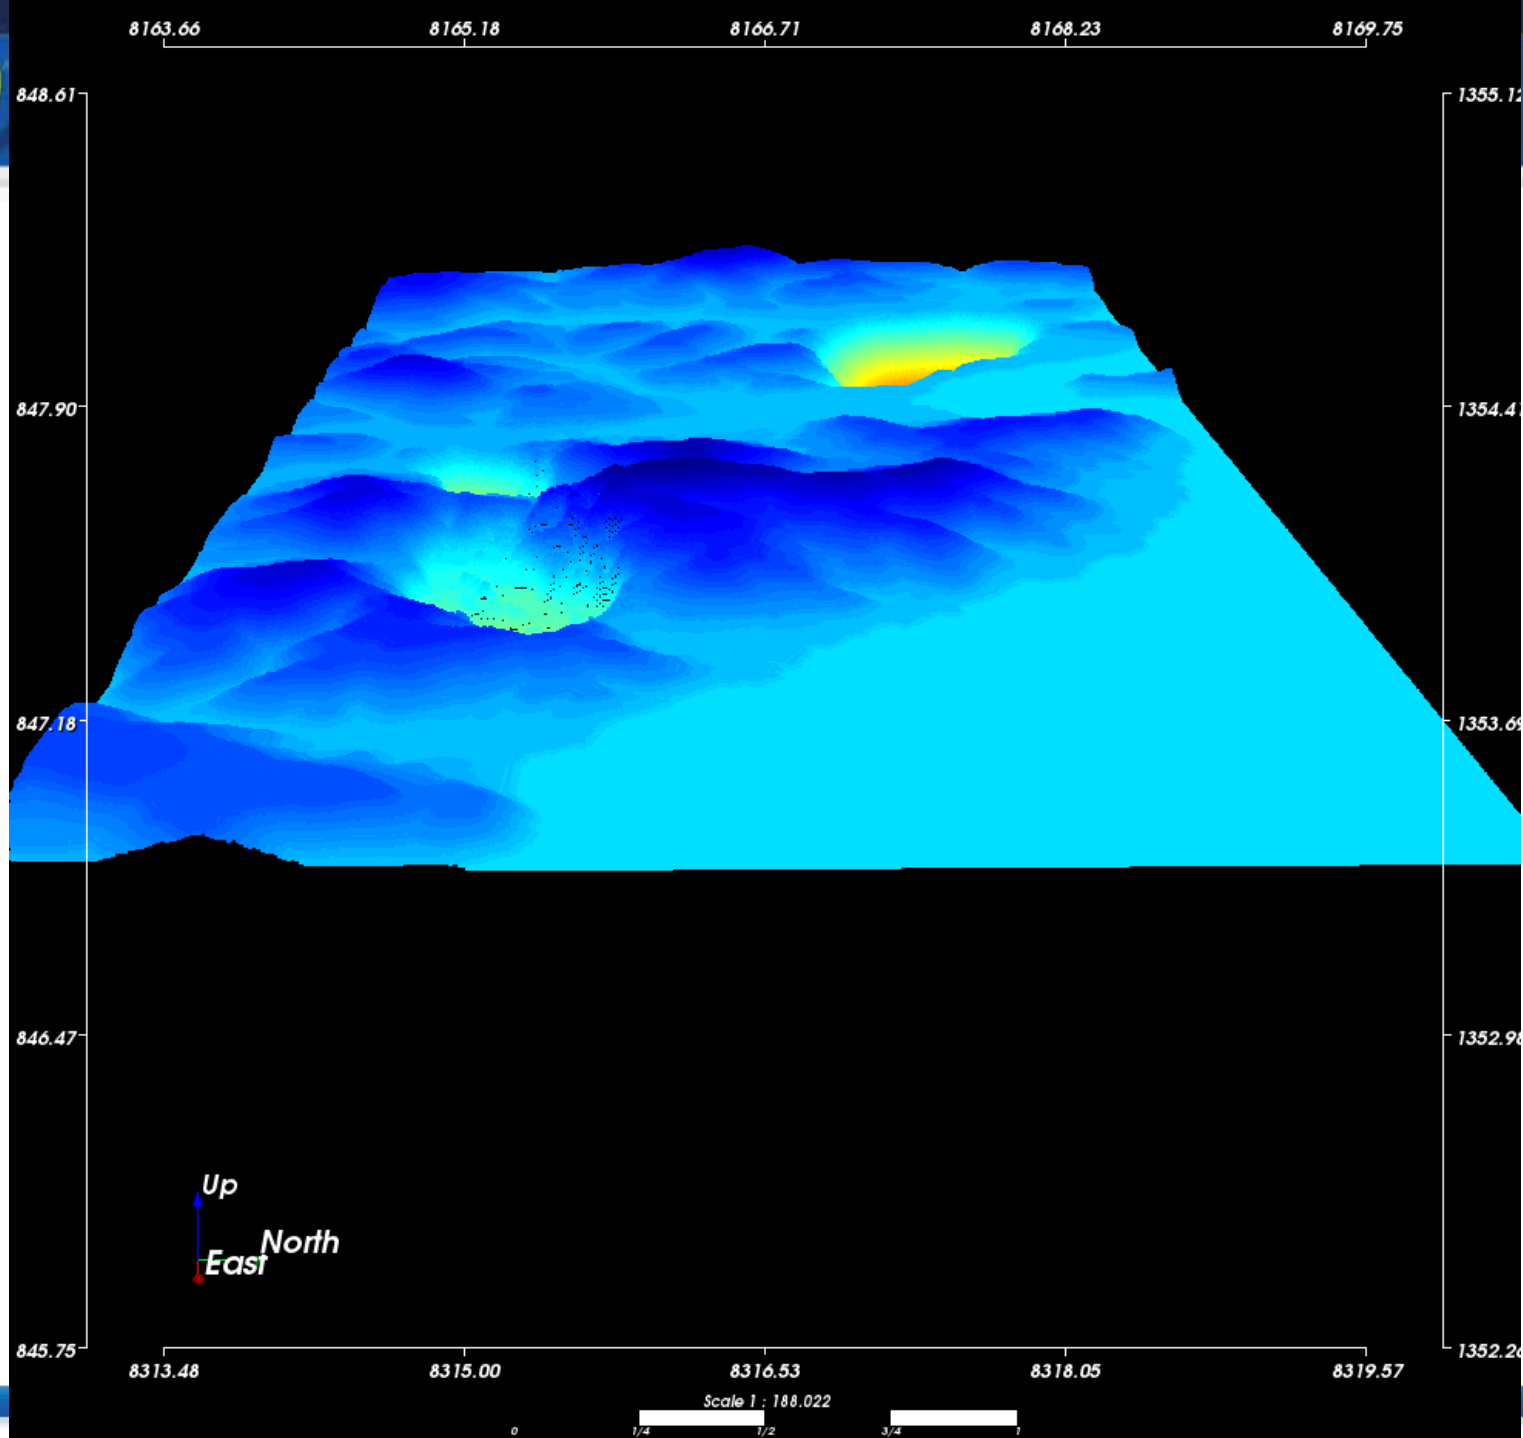







































$$\text{max.} \quad c^T x$$

$$\text{s.t.} \quad Ax \leq b$$


$$x \geq 0$$

# Simplex





## Interior Point





# LINPACK AND MATRIX

Organized around matrix decompositions

LU

Cholesky

QR

singular value

- Small block model  $\sim 66k$  blocks
  - Sparse Cholesky  $23000 \times 23000$  for one iteration
    - 529M elements
    - 207k non zero elements
  - Dense CPU single solution
    - Over 30 min
  - Sparse CPU single solution (not state of the art)
    - Few seconds
- CHOLMOD (suitesparse from NVIDIA)
  - 3x speedup over CPU multithread implementation

$$\begin{aligned} A &= \begin{pmatrix} a_{11} & a_{21} & a_{31} \\ a_{21} & a_{22} & a_{32} \\ a_{31} & a_{32} & a_{33} \end{pmatrix} \\ &= \begin{pmatrix} l_{11} & 0 & 0 \\ l_{21} & l_{22} & 0 \\ l_{31} & l_{32} & l_{33} \end{pmatrix} \begin{pmatrix} l_{11} & l_{21} & l_{31} \\ 0 & l_{22} & l_{32} \\ 0 & 0 & l_{33} \end{pmatrix} \equiv LL^T \\ &= \begin{pmatrix} l_{11}^2 & l_{21}l_{11} & l_{31}l_{11} \\ l_{21}l_{11} & l_{21}^2 + l_{22}^2 & l_{31}l_{21} + l_{32}l_{22} \\ l_{31}l_{11} & l_{31}l_{21} + l_{32}l_{22} & l_{31}^2 + l_{32}^2 + l_{33}^2 \end{pmatrix} \end{aligned}$$

$$l_{kk} = \sqrt{a_{kk} - \sum_{j=1}^{k-1} l_{kj}^2}$$

$$l_{ik} = \frac{1}{l_{kk}} \left( a_{ik} - \sum_{j=1}^{k-1} l_{ij}l_{kj} \right)$$

```
1: procedure CHOLESKY( $A$ )
2:   int  $i, j, k$ ;
3:   for  $k := 0$  to  $n - 1$  do
4:      $A[k, k] := \sqrt{A[k, k]}$ ;    /* Obtain the square root of the diagonal element. */
5:     for  $j := k + 1$  to  $n - 1$  do
6:        $A[k, j] := A[k, j] / A[k, k]$ ;    /* The division step. */
7:     end for
8:     for  $i := k + 1$  to  $n - 1$  do
9:       for  $j := i$  to  $n - 1$  do
10:         $A[i, j] := A[i, j] - A[k, i] \times A[k, j]$ ;    /* The elimination step. */
11:      end for
12:    end for
13:  end for
```

# Single Thread

```
int chol_gold(const Matrix A, Matrix U) {
    unsigned int i, j, k;
    unsigned int size = A.num_rows * A.num_columns;

    // Copy the contents of the A matrix into the working matrix U
    for (i = 0; i < size; i++)
        U.elements[i] = A.elements[i];

    // Perform the Cholesky decomposition in place on the U matrix
    for (k = 0; k < U.num_rows; k++) {
        // Take the square root of the diagonal element
        U.elements[k * U.num_rows + k] = sqrt(U.elements[k * U.num_rows + k]);
        if (U.elements[k * U.num_rows + k] <= 0) {
            printf("Cholesky decomposition failed. \n");
            return 0;
        }

        // Division step
        for (j = (k + 1); j < U.num_rows; j++)
            U.elements[k * U.num_rows + j] /= U.elements[k * U.num_rows + k]; // Division step

        // Elimination step
        for (i = (k + 1); i < U.num_rows; i++)
            for (j = i; j < U.num_rows; j++)
                U.elements[i * U.num_rows + j] -= U.elements[k * U.num_rows + i] * U.elements[k * U.num_rows + j];
    }

    // As the final step, zero out the lower triangular portion of U
    for (i = 0; i < U.num_rows; i++)
        for (j = 0; j < i; j++)
            U.elements[i * U.num_rows + j] = 0.0;
    return 1;
}
```

```

1  #ifndef _MATRIX_H_
2  #define _MATRIX_H_
3
4  // Thread block size
5  #define MATRIX_SIZE 2048
6  #define NUM_THREADS 4
7  #define NUM_OMP_THREADS NUM_THREADS
8
9  // Matrix dimensions
10 #define NUM_COLUMNS MATRIX_SIZE // Number of columns
11 #define NUM_ROWS MATRIX_SIZE // Number of rows
12
13 // Matrix Structure declaration
14 typedef struct {
15     //width of the matrix represented
16     unsigned int num_columns;
17     //height of the matrix represented
18     unsigned int num_rows;
19     //number of elements between the beginning of the rows
20     // rows in the memory layout (useful for column-major access)
21     unsigned int pitch;
22     //Pointer to the first element of the matrix
23     float* elements;
24 } Matrix;
25
26 #endif // _MATRIX_H_

```

**CPU-Z**

CPU Caches Mainboard Memory SPD Graphics About

**Processor**

Name	Intel Core i5 3330		
Code Name	Ivy Bridge	Max TDP	77 W
Package	Socket 1155 LGA		
Technology	22 nm	Core Voltage	0.808 V

Specification: Intel(R) Core(TM) i5-3330 CPU @ 3.00GHz

Family	6	Model	A	Stepping	9
Ext. Family	6	Ext. Model	3A	Revision	E1
Instructions	MMX, SSE (1, 2, 3, 3S, 4.1, 4.2), EM64T, VT-x, AES, AVX				

**Clocks (Core #0)**

Core Speed	3109.31 MHz
Multiplier	x 31.0
Bus Speed	100.3 MHz
Rated FSB	

**Cache**

L1 Data	4 x 32 KBytes	8-way
L1 Inst.	4 x 32 KBytes	8-way
Level 2	4 x 256 KBytes	8-way
Level 3	6 MBytes	12-way

Selection: Processor #1 Cores: 4 Threads: 4

**CPU-Z** Version 1.61.5.x64 Validate OK

```

// Perform the Cholesky decomposition in place on the U matrix
for(k = 0; k < U.num_rows; k++)
{
    //Only one thread does square root and division
    if(id==0)
    {
        // Take the square root of the diagonal element
        U.elements[k * U.num_rows + k] = sqrt(U.elements[k * U.num_rows + k]);
        if(U.elements[k * U.num_rows + k] <= 0){
            printf("Cholesky decomposition failed. \n");
            return 0;
        }

        // Division step
        for(j = (k + 1); j < U.num_rows; j++)
        {
            U.elements[k * U.num_rows + j] /= U.elements[k * U.num_rows + k]; // Division step
        }
    }

    //Sync threads!!!!
    sync_pthreads(barrier, id);

    //For this k iteration, split up i
    //Size of i range originally
    int isize = U.num_rows - (k + 1);
    int items_per_thread, items_last_thread;
    range_splitter(isize, NUM_PTHREADS, &items_per_thread, &items_last_thread);
    //Divy up work
    //Elim work
    int elimi_start, elimi_end;
    int offset = (k + 1); //To account for not starting at i=0 each time

```



```

// Perform the Cholesky decomposition in place on the U matrix
for(k = 0; k < U.num_rows; k++)
{
    //Only one thread does square root and division
    if(id==0)
    {
        // Take the square root of the diagonal element
        U.elements[k * U.num_rows + k] = sqrt(U.elements[k * U.num_rows + k]);
        if(U.elements[k * U.num_rows + k] <= 0){
            printf("Cholesky decomposition failed. \n");
            return 0;
        }

        // Division step
        for(j = (k + 1); j < U.num_rows; j++)
        {
            U.elements[k * U.num_rows + j] /= U.elements[k * U.num_rows + k]; // Division step
        }
    }

    //Sync threads!!!!
    sync_pthreads(barrier, id);

    //For this k iteration, split up i
    //Size of i range originally
    int isize = U.num_rows - (k + 1);
    int items_per_thread, items_last_thread;
    range_splitter(isize, NUM_PTHREADS, &items_per_thread, &items_last_thread);
    //Divy up work
    //Elim work
    int elimi_start, elimi_end;
    int offset = (k + 1); //To account for not starting at i=0 each time

```

```

// Perform the Cholesky decomposition in place on the U matrix
for(k = 0; k < U.num_rows; k++)
{
    //Only one thread does square root and division
    if(id==0)
    {
        // Take the square root of the diagonal element
        U.elements[k * U.num_rows + k] = sqrt(U.elements[k * U.num_rows + k]);
        if(U.elements[k * U.num_rows + k] <= 0){
            printf("Cholesky decomposition failed. \n");
            return 0;
        }


        // Division step
        for(j = (k + 1); j < U.num_rows; j++)
        {
            U.elements[k * U.num_rows + j] /= U.elements[k * U.num_rows + k]; // Division step
        }
    }

    //Sync threads!!!!
    sync_pthreads(barrier, id);

    //For this k iteration, split up i
    //Size of i range originally
    int isize = U.num_rows - (k + 1);
    int items_per_thread, items_last_thread;
    range_splitter(isize, NUM_PTHREADS, &items_per_thread, &items_last_thread);
    //Divy up work
    //Elim work
    int elimi_start, elimi_end;
    int offset = (k + 1); //To account for not starting at i=0 each time

```

## Speedup vs. Number of Threads


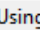
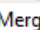
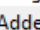







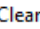


# Pthread: our optimizations

Current branch

SHOW REMOTE BRANCHES

Date Order

Graph	Description
 origin/master	Better timing for cpu single and pthreads
 origin/HEAD	Using only line to detect pthread library
 master	Merge pull request #1 from alexbenfica/master
	Added .bat to make build easier on Windows.
	Simplifying index calculations
	Fix chol pthread bug buffer overflow
	Fixing linux compilation
	Remove openmp code
	Creating Makefile to compile under windows and timing support
	Change number of threads to 8
	Cleanup end spaces
	All files

Sorted by file status

Commit: 105f473d409157ebee7c256bfc9388c4538c1e83 [105f473]

Parents: 52977e4aad

Author: Fabrício Ceolin <fabricio.ceolin@miningmath.com>

Date: terça-feira, 25 de novembro de 2014 11:14:36

Simplifying index calculations

chol.c

chol.h

Hunk 1 : Lines 122-137

```

122 122 {
123 123     //Divide up total size by number of threads
124 124     //How many are left over?
125 125     int elems_left_over = size%num_threads;
126 126     int elements_per_thread = size/num_threads;
127 127     int last_thread_elements = elements_per_thread;
128 128     +
128 129     if(elems_left_over !=0)
129 130     {
130 131         -         int not_total = elements_per_thread*(num_threads-1);
131 132         -         last_thread_elements = size - not_total;
132 133         -
133 134         -         if(last_thread_elements<0)
134 135         -         {
135 136         -             //Too much now
136 137         -             elements_per_thread-=2;
137 138         -             not_total = elements_per_thread*(num_threads-1);
138 139         -             last_thread_elements = size - not_total;
139 139         -         }
140 140         +         last_thread_elements = elements_per_thread+elems_left_over;
141 141     }
142 142     //Double check because math is hard
143 143     if( (((num_threads-1)*elements_per_thread) + last_thread_elements) !=
144 144     {
145 145         printf("AH! MATH! threads:%d elementsperthread:%d lastthreadelm:%d

```

```
D:\Dropbox\Devel\Win\CudaProject\Drexel-ECE622-Midterm1\build>CholThread.exe
Creating a 2048 x 2048 matrix with random numbers between [-.5, .5]...done.
Generating the symmetric matrix...done.
Generating the positive definite matrix...done.
Performing Cholesky decomposition on the CPU using the single-threaded version.

Run time:      9.4130001068 s.
Cholesky decomposition on the CPU was successful.

Performing Cholesky decomposition on the CPU using the PTHREAD version.
Run time:      3.0450000763 s.
Speedup from single to pthreads = 3.091297 x.
Double checking for correctness by recovering the original matrix.
```

# CUDA device query: Geforce GTS 450

```
CUDA Device Query (Runtime API) version (CUDART static linking)

Detected 1 CUDA Capable device(s)

Device 0: "GeForce GTS 450"
  CUDA Driver Version / Runtime Version      6.5 / 6.5
  CUDA Capability Major/Minor version number: 2.1
  Total amount of global memory:             1024 MBytes (1073741824 bytes)
  ( 4) Multiprocessors, ( 48) CUDA Cores/MP: 192 CUDA Cores
  GPU Clock rate:                           1620 MHz (1.62 GHz)
  Memory Clock rate:                         1804 Mhz
  Memory Bus Width:                          128-bit
  L2 Cache Size:                             262144 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(65536), 2D=(65536, 65535), 3D=(2048, 2048, 2048)
  Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total number of registers available per block: 32768
  Warp size:                                  32
  Maximum number of threads per multiprocessor: 1536
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (65535, 65535, 65535)
  Maximum memory pitch:                       2147483647 bytes
  Texture alignment:                          512 bytes
  Concurrent copy and kernel execution:       Yes with 1 copy engine(s)
  Run time limit on kernels:                   Yes
  Integrated GPU sharing Host Memory:          No
  Support host page-locked memory mapping:    Yes
  Alignment requirement for Surfaces:         Yes
  Device has ECC support:                     Disabled
  CUDA Device Driver Mode (TCC or WDDM):       WDDM (Windows Display Driver Model)
  Device supports Unified Addressing (UVA):    Yes
  Device PCI Bus ID / PCI location ID:        1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >

deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 6.5, CUDA Runtime Version = 6.5, NumDevs = 1, Device0 = GeForce GTS 450
Result = PASS
```

# CUDA device query: Geforce GTS 450

```
CUDA Device Query (Runtime API) version (CUDART static linking)
```

```
Detected 1 CUDA Capable device(s)
```

```
Device 0: "GeForce GTS 450"
  CUDA Driver Version / Runtime Version      6.5 / 6.5
  CUDA Capability Major/Minor version number: 2.1
  Total amount of global memory:             1024 MBytes (1073741824 bytes)
  ( 4) Multiprocessors, ( 48) CUDA Cores/MP: 192 CUDA Cores
  GPU Clock rate:                            1620 MHz (1.62 GHz)
  Memory Clock rate:                         1804 Mhz
  Memory Bus Width:                          128-bit
  L2 Cache Size:                             262144 bytes
  Maximum Texture Dimension Size (x,y,z)     1D=(65536), 2D=(65536, 65535), 3D=(2048, 2048, 2048)
  Maximum Layered 1D Texture Size, (num) layers 1D=(16384), 2048 layers
  Maximum Layered 2D Texture Size, (num) layers 2D=(16384, 16384), 2048 layers
  Total amount of constant memory:            65536 bytes
  Total amount of shared memory per block:    49152 bytes
  Total number of registers available per block: 32768
  Warp size:                                  32
  Maximum number of threads per multiprocessor: 1536
  Maximum number of threads per block:        1024
  Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
  Max dimension size of a grid size (x,y,z):  (65535, 65535, 65535)
  Maximum memory pitch:                       2147483647 bytes
  Texture alignment:                          512 bytes
  Concurrent copy and kernel execution:       Yes with 1 copy engine(s)
  Run time limit on kernels:                   Yes
  Integrated GPU sharing Host Memory:          No
  Support host page-locked memory mapping:    Yes
  Alignment requirement for Surfaces:          Yes
  Device has ECC support:                      Disabled
  CUDA Device Driver Mode (TCC or WDDM):       WDDM (Windows Display Driver Model)
  Device supports Unified Addressing (UVA):    Yes
  Device PCI Bus ID / PCI location ID:        1 / 0
  Compute Mode:
    < Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
```

```
deviceQuery, CUDA Driver = CUDART, CUDA Driver Version = 6.5, CUDA Runtime Version = 6.5, NumDevs = 1, Device0 = GeForce GTS 450
Result = PASS
```



```
int threads_per_block_sqrt = 512;
int blocks_sqrt = MATRIX_SIZE / threads_per_block_sqrt;
dim3 thread_block(threads_per_block_sqrt, 1, 1);
dim3 grid(blocks_sqrt, 1);
chol_kernel_cudaUFMG_sqrt <<<grid, thread_block>>>(gpu_u.elements);
```



```
__global__ void chol_kernel_cudaUFMG_sqrt(float * U) {  
    // Get a thread identifier  
    int tx = blockIdx.x * blockDim.x + threadIdx.x;  
    int tx_diag = tx * MATRIX_SIZE + tx;  
    U[tx_diag] = sqrt(U[tx_diag]);  
}
```



```
int block_x_div = 16;
int block_y_div = 16;
int thread_x_div = 4;
int thread_y_div = 4;
dim3 grid_div(block_x_div, block_y_div, 1);
dim3 thread_block_div(thread_x_div, thread_y_div, 1);
int elements_per_thread_div = ((MATRIX_SIZE * MATRIX_SIZE) / 2) / (thread_x_div * thread_y_div * block_x_div * block_y_div);
chol_kernel_cudaUFMG_division <<<grid_div, thread_block_div >>>(gpu_u.elements, elements_per_thread_div);
```

# CUDA div kernel

```

__global__ void chol_kernel_cudaUFMG_division(float * U, int elem_per_thr) {
    // Get a thread identifier
    int tx = blockIdx.x * blockDim.x + threadIdx.x;
    int ty = blockIdx.y * blockDim.y + threadIdx.y;

    int tn = ty * blockDim.x * gridDim.x + tx;

    for(unsigned i=0;i<elem_per_thr;i++){
        int iel = tn * elem_per_thr + i;
        int xval = iel % MATRIX_SIZE;
        int yval = iel / MATRIX_SIZE;

        if(xval == yval){
            continue;
        }

        // if on the lower diagonal...
        if(yval > xval){
            xval = MATRIX_SIZE - xval - 1;
            yval = MATRIX_SIZE - yval - 1;
        }

        int iU = xval + yval * MATRIX_SIZE;
        int iDiag = yval + yval * MATRIX_SIZE;

        U[iU] /= U[iDiag];
    }
}

```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
1	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
2	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
3	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
4	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
5	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
6	Red	Red	Red	Red	Red	Red	Blue	Green	Green	Green	Green	Green	Green	Green	Green	Green
7	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
8	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
9	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Green	Green	Green	Green	Green	Green	Green
10	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
11	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow
12	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow
13	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow
14	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow
15	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow

# CUDA elimination

```
int block_y_eli = 1;
//Each thread within a block will take some j iterations
int thread_x_eli = 256;
int thread_y_eli = 1;


//Each kernel call will be one iteration of out K loop
for (int k = 0; k < MATRIX_SIZE; k++) {

    //Want threads to stride across memory
    //i is outer loop
    //j is inner loop
    //so threads should split the j loop
    //Each thread block will take an i iteration

    // i=k+1;i<MATRIX_SIZE
    int isize = MATRIX_SIZE - (k + 1);
    if(isize==0){
        isize++;
    }
    int block_x_eli = isize;

    //Set up the execution grid on the GPU
    dim3 thread_block(thread_x_eli, 1, 1);
    dim3 grid(block_x_eli, 1);

    //Call kernel with for this K iteration
    chol_kernel_cudaUFMG_elimination <<<grid, thread_block>>>(gpu_u.elements, k);
}
```



```
chol_kernel_cudaUFMG_zero <<<grid_div, thread_block_div>>>(gpu_u.elements, elements_per_thread_div);
```

# CUDA elimination kernel

```
__global__ void chol_kernel_cudaUFMG_elimination(float * U, int k) {
```

```
    //This call acts as a single K iteration
```

```
    //Each block does a single i iteration
```

```
    int i = (k+1) + blockIdx.x;
```

```
    //Each thread does some part of j
```

```
    //Stride in units of 'stride'
```

```
    //Thread 0 does 0, 16, 32
```

```
    //Thread 1 does 1, 17, 33
```

```
    //..etc.
```

```
    int jstart = i + threadIdx.x;
```

```
    int jstep = blockDim.x;
```

```
    // Pre-calculate indexes outside loop
```

```
    int km = k * MATRIX_SIZE;
```

```
    int im = i * MATRIX_SIZE;
```

```
    int ki = km + i;
```



```
    //Do work for this i iteration
```

```
    for (int j=jstart; j<MATRIX_SIZE; j+=jstep) {
```

```
        U[im + j] -= U[ki] * U[km + j];
```

```
    }
```

```
}
```



```

__global__ void chol_kernel_cudaUFMG_zero(float * U, int elem_per_thr) {
    // Get a thread identifier
    int tx = blockIdx.x * blockDim.x + threadIdx.x;
    int ty = blockIdx.y * blockDim.y + threadIdx.y;

    int tn = ty * blockDim.x * gridDim.x + tx;

    for(unsigned i=0;i<elem_per_thr;i++){
        int iel = tn * elem_per_thr + i;
        int xval = iel % MATRIX_SIZE;
        int yval = iel / MATRIX_SIZE;

        if(xval == yval){
            continue;
        }

        // if on the upper diagonal...
        if(yval < xval){
            xval = MATRIX_SIZE - xval - 1;
            yval = MATRIX_SIZE - yval - 1;
        }

        int iU = xval + yval * MATRIX_SIZE;
        U[iU] = 0;
    }
}

```

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
1	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
2	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
3	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
4	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
5	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
6	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
7	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
8	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
9	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow
10	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow	Yellow
11	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow	Yellow
12	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow	Yellow
13	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow	Yellow
14	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue	Yellow
15	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Yellow	Blue



## SPEEDUP Results

TYPE	SINGLE	PTHREAD	OUR PTHREAD	CUDA	OUR CUDA
SINGLE	1.00				
PTHREAD	2.20	1.00			
OUR PTHREAD	3.10	<b>1.41</b>	1.00		
CUDA	24.00	10.91	7.74	1.00	
OUR CUDA	33.00	15.00	10.65	<b>1.38</b>	1.00

- Improvements
  - It is possible to speedup using our algorithm with CUDA
  - No room for shared memory in our implementation
    - (We've tried other methods, without success)



- Future work
  - Pthread
    - Split sqrt and division over threads
  - Study Sparse Matrix algorithm
    - suitesparse with metis
    - Test CUDA and CPU multithread

# Questions?



**Fabrício Ceolin e Alex Teixeira**